

Product Build & Investment Report

AI Revenue Recovery OS — Technology · Security · Cost · AI Workforce

“Predict. Prevent. Work. Resolve. Recover.” — an AI orchestration & intelligence layer for healthcare revenue.

PREPARED FOR	AI Revenue Recovery OS
DATE	August 18, 2026
MVP INVESTMENT (est.)	\$383,068 (Rs 3.45 Cr)
12 / 24-MONTH (est.)	\$1,492,847 / \$4,450,298
CLASSIFICATION	Confidential · Investor & engineering · MODEL ASSUMPTIONS

Disclosures

Prices are researched 2026 figures with sources; usage-derived monthly figures are labeled MODEL ASSUMPTION and are configurable in the workbook. Quote-only prices (ChatGPT Enterprise, some Azure dedicated capacity) are recorded 'Not publicly disclosed' — never invented. “HIPAA-aligned architecture / HIPAA program” — never “HIPAA certified”; Azure BAA availability does not by itself make the application compliant. Development uses synthetic/de-identified data only. Payer-side is AI decision-support with human/payer governance. No PHI.

Contents

- 1. Executive Summary 4
- 2. Product Vision & Models 5
- 3. Functional Scope 6
- 4. Technology Stack & Architecture 7
- 5. Multi-Tenant Architecture 8
- 6. AI Architecture & Agent Workforce 9
- 7. Security, HIPAA & Compliance 10
- 8. Azure Cost 11
- 9. AI Cost — Runtime & Developer Tools 12
 - 9.1 Claude Max (developer productivity) 12
 - 9.2 ChatGPT 13
- 10. Human Resources & Development Cost 14
- 11. Milestones & Project Plan 15
- 12. Investment & Budgets 16
- 13. Unit Economics 17
- 14. Provider & Payer Architecture 18
 - 14.1 Provider architecture 18
 - 14.2 Payer architecture (decision-support) 18
- 15. AI Agent Workforce 19
- 16. Azure Cost by Scale 20
 - 16.1 Medium Production — \$7,176/month 20
 - 16.2 Large Production — \$38,694/month 20
- 17. Human Resources — Salary Bands 21
- 18. Milestones — Full Detail 22
- 19. 12-Month & 24-Month Budgets 23
 - 19.1 12-month budget 23
 - 19.2 24-month budget 23
- 20. Commercial Cost Models 24
 - 20.1 Build / Buy / Partner 24
- 21. Disaster Recovery, Continuity & Governance 25

21.1 RPO / RTO	25
21.2 Continuity responses	25
21.3 AI agent governance	25
22. Connectivity Strategy	26
23. Security Controls & PHI Data Flow	27
23.1 PHI data flow (AI Gateway)	27
24. Risks	28
25. Final Recommendation	28
Appendix. Pricing Sources	29

1. Executive Summary

AI Revenue Recovery OS is an AI-native orchestration and intelligence layer for U.S. healthcare revenue, spanning provider/RCM and (later) payer workflows across EDI → API → FHIR → payer portal → AI voice → human. This report models the technology, security, team and cost to build it — MVP through 24 months — on Microsoft Azure with an AI-augmented engineering organization.




Figure 1. Estimated total investment: MVP, 12-month and 24-month (MODEL ASSUMPTION).

Headline	USD	INR (Rs , FX 90)
MVP investment	\$383,068	Rs 34,476,120
12-month investment	\$1,492,847	Rs 134,356,230
24-month investment	\$4,450,298	Rs 400,526,820
Avg monthly burn (yr 1)	\$124,404	Rs 11,196,360
Azure — MVP / month	\$1,377	Rs 123,930
Claude Max baseline / month (5×20x)	\$1,000	Rs 90,000
All-in cost per worked claim	\$0.43	Rs 39

Recommendation

Build the common platform from day one; **launch Provider/RCM first**; add Payer (decision-support) after validation. Baseline team of **5 FTE × Claude Max 20x**; Azure Container Apps + PostgreSQL Flexible Server for the MVP, private networking + Defender/Sentinel at production; HIPAA-aligned with SOC 2 on the roadmap. Do not over-engineer the MVP.

2. Product Vision & Models

The platform determines what needs to happen, why, through which channel, whether it can be solved digitally, and — only when digital fails — whether a portal, AI voice or human should act; then captures the outcome and the revenue protected or recovered. It is positioned as an orchestration and intelligence layer, not a chatbot, caller, billing app or dashboard.

Model	Customer	Positioning
A — Provider / RCM	Billing/RCM, hospitals, health systems, groups	Provider revenue automation across the RCM lifecycle
B — Payer	Insurers, health plans, MA orgs, Medicaid MCOs, TPAs	AI decision-support + workflow automation with governance
C — Provider + Payer	Long-term	Shared foundation; strict tenant isolation

Payer-side guardrail


Final coverage / medical-necessity / payment / adjudication decisions remain with the payer's qualified personnel. AI cannot be the sole basis for a medical-necessity decision (CMS-4201-F; California SB 1120). Delegated payer operations require separate regulatory/legal review.

3. Functional Scope

- **Provider:** Eligibility · Benefits · Patient Responsibility Estimation · Prior Auth · Claims · Claim Status · Denial Intelligence · Appeals · Revenue Recovery · Payer Intelligence.
- **Payer:** Claim Intake · Validation · Eligibility/Benefit Validation · Adjudication Decision-Support · Payment Integrity · Denial Ops · Appeal Admin · Provider Communication · Payer Intelligence · Analytics.
- **Platform:** AI Workforce · Human Exception Center · Analytics · Audit · Security · Administration.
- **Standards:** X12 270/271, 276/277, 277CA, 278, 837P/I/D, 835, 999, TA1, 275; FHIR R4 (Coverage, Claim, ClaimResponse, EOB, CoverageEligibility*, Task, DocumentReference, Questionnaire); HL7 Da Vinci; CMS interoperability (CMS-0057-F FHIR APIs by 2027 for CMS-regulated payers).

4. Technology Stack & Architecture

Microsoft Azure is the primary cloud. The MVP is deliberately lean; premium networking and security are added at production.



Security plane: Entra · Key Vault · Private Link · Firewall · Defender · Sentinel · Monitor · Backup/DR

Figure 2. Conceptual architecture — secured ingress, provider/payer/admin APIs, AI orchestration (rules + agents + RAG), workflow engine, tenant-isolated data, and a full security plane.

Layer	Technology
Compute	Azure Container Apps (AKS when justified)
Data	PostgreSQL Flexible Server (RLS, HA, PITR); Redis
Messaging/Storage	Service Bus; Blob (private, encrypted)
AI	AI Gateway → Azure OpenAI (PHI/BAA) · Anthropic · OpenAI; Azure AI Search (RAG)
Edge/API	Front Door + WAF + API Management; Entra ID
Security	Key Vault/HSM, Private Link, Firewall, Defender, Sentinel, Monitor

5. Multi-Tenant Architecture

Tenant types: PLATFORM · PROVIDER · PAYER · RCM/BPO · ENTERPRISE. Every record carries tenant_id, organization_id, environment_id. Provider and payer data never automatically cross; any permitted exchange is authorized, purpose-limited, logged, encrypted and tenant-aware.

Stage	Isolation model
MVP	Shared DB + Row-Level Security + app authorization + RBAC + encryption + audit
V1	+ schema separation for large tenants; ABAC; per-tenant keys where justified
Enterprise	Database-per-tenant / dedicated environment, keys, networking, logging; optional regional deployment

6. AI Architecture & Agent Workforce

A provider-agnostic AI Gateway routes by task, cost, latency, reasoning need, PHI policy and confidence: small/fast models for classification/extraction, strong models for reasoning/appeals/exceptions, with a fallback chain of primary → secondary → rules → human. A secure, tenant-filtered RAG pipeline grounds answers in payer/CMS/medical policies with citations and confidence.

Twenty specialized AI agents (Product/Project Manager, Business Analyst, Healthcare RCM, Payer Research, Architecture, Backend/Frontend/AI Dev, EDI, FHIR, Integration, Database, QA, Security, DevOps, FinOps, Documentation, Compliance Research, Test Data) are operated by the human team through AI tools — **not** one paid subscription per agent. Each agent has scoped tool/data/tenant/action permissions and full audit; high-risk actions require human approval.

PHI AI security

AI Gateway: PHI detection → policy engine → redaction/tokenization → approved model (Azure OpenAI under BAA for the PHI path) → response validation → audit. Development uses synthetic/de-identified data only; developer AI tools (Claude Max, ChatGPT) are never production PHI infrastructure.

7. Security, HIPAA & Compliance

Zero-Trust, least-privilege, defense-in-depth. Entra ID (MFA/CA/PIM/JIT), Key Vault/CMK/HSM, Private Link, Firewall, WAF, DDoS, Defender for Cloud, Sentinel SIEM, immutable audit, encrypted backup + DR. HIPAA is an ongoing program (risk assessment, policies/SOPs, safeguards, training, evidence) — the Azure BAA is included via Microsoft Product Terms/DPA but does not by itself confer compliance.




Figure 3. Security & compliance cost items (benchmark ranges → mid).

Framework	Timing	Cost range	Business value
SOC 2 Type I	M10–M11	\$15k–35k	Provider trust; sales enabler
SOC 2 Type II	M12 + 6-mo window	\$30k–80k	Enterprise/payer requirement
HITRUST i1	Year 2	\$60k–200k	Payer contracts; high assurance
HIPAA program	M10 ongoing	\$20k–40k + tooling	Table stakes for PHI

8. Azure Cost

Azure monthly cost scales from a lean MVP to enterprise production. Figures are MODEL ASSUMPTIONS built from published unit rates; the workbook lets every line and the USD/INR rate be edited.




Figure 4. Azure monthly cost by scale.

Service (MVP)	Tier	Monthly USD
Container Apps	Consumption (scaled)	\$120
PostgreSQL Flexible Server	GP D4ds_v5 + HA	\$500
Cache for Redis	Standard C1	\$52
Blob Storage	Hot LRS	\$40
Service Bus	Standard	\$30
Key Vault	Standard	\$15
Front Door	Standard + custom WAF	\$35
API Management	Basic v2	\$210
Defender for Cloud	Servers P2 (few)	\$60
Sentinel + Log Analytics	~30 GB/mo	\$150

MVP Azure total ≈ \$1,377/mo; Medium ≈ \$7,176/mo; Large ≈ \$38,694/mo. Full line items by scale are in the workbook (AZURE COST BY SCALE).

9. AI Cost — Runtime & Developer Tools

Two separate budgets: (a) production AI runtime (per-token/per-minute), and (b) AI developer subscriptions. A ChatGPT/Claude subscription is never production API credit.




Figure 5. Cost per worked claim (AI + voice + document AI + clearinghouse transaction).




Figure 6. Production AI runtime cost by scale.

9.1 Claude Max (developer productivity)

Claude Max 5x is \$100/user/mo and Max 20x is \$200/user/mo (per individual; Claude Code included). Baseline: 5 users $\times$ Max 20x = \$1,000/mo (\$12,000/yr). Do not create one subscription per AI agent.

Users	Plan	Monthly	Annual
3	Max 5x	\$300	\$3,600
3	Max 20x	\$600	\$7,200
5	Max 5x	\$500	\$6,000
5	Max 20x	\$1,000	\$12,000
10	Max 5x	\$1,000	\$12,000
10	Max 20x	\$2,000	\$24,000
15	Max 5x	\$1,500	\$18,000
15	Max 20x	\$3,000	\$36,000

9.2 ChatGPT

ChatGPT Business (formerly Team) is \$20/user/mo (annual, 2-seat min); Enterprise is quote-only (not publicly disclosed). Model separately from the OpenAI/Azure OpenAI production API.

10. Human Resources & Development Cost

An India-based engineering team (US customers) with an AI productivity uplift. Salary bands are benchmark ranges (Glassdoor/6figr/Levels), fully-loaded at ×1.4.




Figure 7. Annual human-resource cost by team model (mid-band, loaded).

Team model	Roles	Annual (mid, loaded)
Lean MVP (5)	5	\$310,800
Balanced MVP (8)	8	\$414,400
Production (13)	13	\$624,400

MVP build (lean team × 8 mo) ≈ \$207,200. AI productivity scenarios of +10/20/30/40% are modeled; plan on +30%.

11. Milestones & Project Plan




Figure 8. M0–M12 milestone timeline (~46 weeks realistic).

Scenario	MVP (M9)	Pilot (M11)	Production (M12)
Aggressive (0.8×)	~26 wk	~30 wk	~37 wk
Realistic (1.0×)	~32 wk	~40 wk	~46 wk
Conservative (1.3×)	~42 wk	~52 wk	~60 wk

PM platform: GitHub Projects (recommended). Workflow: Requirement → AI PM → story → AI architect → AI dev → AI QA/Security → PR → human review → CI/CD → staging → human approval → production. Protected main; mandatory PR + SAST/DAST/SCA/secret/container/IaC scanning.

12. Investment & Budgets




Figure 9. MVP investment breakdown.




Figure 10. Cumulative investment over 24 months.

Phase	Investment (USD)	INR (Rs)
MVP (one-time, ~8 mo)	\$383,068	Rs 34,476,120
12-month	\$1,492,847	Rs 134,356,230
24-month	\$4,450,298	Rs 400,526,820

Contingency scenarios of 10% / 15% / 20% are modeled (15% used above). Every cost is classified one-time / monthly / annual / usage / fixed / variable / estimated / publicly-priced / quote-required in the workbook.

13. Unit Economics

Blended all-in cost per worked claim ≈ \$0.43 (AI \$0.12 + voice \$0.06 + document AI \$0.01 + clearinghouse \$0.24). The clearinghouse transaction is the largest per-claim variable cost; Stedi's free tier covers the MVP, and model routing + prompt caching + batch keep AI cost low.

Action	AI \$	Channel \$	Total \$
Eligibility check	0.005	0.12	0.125
Claim status	0.005	0.12	0.125
Prior authorization	0.05	0.12	0.17
Denial analysis	0.10	0.00	0.10
Appeal draft	0.15	0.00	0.15
AI agent action	0.02	0.00	0.02

14. Provider & Payer Architecture

14.1 Provider architecture

Layer	Description
Ingress	Front Door → WAF → API Management → Entra ID
Provider API	Eligibility · Benefits · Estimation · Prior Auth · Claims · Status · Denials · Appeals · Recovery
Orchestration	Right-channel engine: EDI → API → FHIR → Portal → Voice → Human
Data	PostgreSQL (RLS by tenant) · Blob (claims/docs) · Redis
Connectivity	Clearinghouse APIs (Stedi/Optum/Availity); payer FHIR (2027)

14.2 Payer architecture (decision-support)

Layer	Description
Payer API	Claim intake · validation · adjudication decision-support · payment integrity · denial ops · appeals · provider comms
Governance	AI recommendation → rules validation → human/payer approval → execution → audit
Guardrail	No final coverage/medical-necessity/payment decision by AI alone
Data	Payer tenant DB / dedicated for enterprise; never visible to provider tenants
Intelligence	Payer Intelligence + coverage-policy RAG (free CMS NCD/LCD; license InterQual/MCG)

15. AI Agent Workforce

Twenty specialized agents operated by the human team through AI tools — not one subscription per agent. Each has scoped tool/data/tenant/action permissions and full audit; high-risk actions are human-approved.

#	Agent	#	Agent
1	Product Manager	11	FHIR
2	Project Manager	12	Integration
3	Business Analyst	13	Database
4	Healthcare RCM	14	QA
5	Payer Research	15	Security
6	Architecture	16	DevOps
7	Backend Developer	17	FinOps
8	Frontend Developer	18	Documentation
9	AI Engineer	19	Compliance Research
10	EDI	20	Test Data

16. Azure Cost by Scale

Full monthly line items per scenario (MODEL ASSUMPTIONS from published unit rates). The workbook exposes every cell plus the USD/INR rate.

16.1 Medium Production — \$7,176/month

Service	Tier	Monthly USD
Container Apps	Consumption (autoscale)	\$600
PostgreSQL Flexible Server	GP large + HA + replica	\$1,400
Cache for Redis	Standard C2/C3	\$250
Blob Storage	Hot + lifecycle	\$200
Service Bus	Premium 1 MU	\$700
Key Vault	Premium (HSM-keys)	\$50
Front Door	Premium (WAF+PL) + traffic	\$530
API Management	Standard v2	\$734
Defender for Cloud	Servers/DB/Containers	\$300
Sentinel + Log Analytics	~150 GB/mo	\$700
Azure Monitor / App Insights	PAYG	\$150
Azure Firewall	Standard	\$912
Azure Bastion	Standard	\$150
Container Registry	Premium	\$50
Bandwidth / Egress	~2 TB	\$200
Azure AI Search	S1	\$250

16.2 Large Production — \$38,694/month

Service	Tier	Monthly USD
Compute (Container Apps / AKS)	Enterprise autoscale	\$3,000
PostgreSQL Flexible Server	Large + HA + replicas + geo	\$6,000
Cache for Redis	Enterprise / Managed Redis	\$1,500
Blob Storage	Hot + GRS + geo	\$1,200
Service Bus	Premium 4 MU	\$2,700
Key Vault + Managed HSM	Premium + HSM pool	\$2,200
Front Door	Premium + high traffic	\$2,000
API Management	Premium (multi-region)	\$6,000
Defender for Cloud	Broad coverage	\$1,500
Sentinel + Log Analytics	~500 GB/mo	\$2,500
Azure Monitor / App Insights	PAYG	\$600
Azure Firewall	Premium + DR region	\$2,600
Azure Bastion	Standard (multi)	\$300
Container Registry	Premium geo-replicated	\$150
Bandwidth / Egress	~15 TB	\$1,500
Azure AI Search	S2/S3	\$2,000
DDoS Network Protection	Standard	\$2,944

17. Human Resources — Salary Bands

India annual base salary bands (benchmark ranges; fully-loaded at ×1.4). Ranges, not quotes.

Role	Low \$	Mid \$	High \$	Loaded Mid \$
CTO / Principal Architect	\$46,000	\$86,000	\$172,000	\$120,400
Product Manager	\$23,000	\$40,000	\$63,000	\$56,000
Healthcare / RCM Business Analyst	\$7,000	\$14,000	\$23,000	\$19,600
Senior Backend Engineer	\$21,000	\$34,000	\$57,000	\$47,600
Senior Frontend Engineer	\$17,000	\$30,000	\$52,000	\$42,000
AI / ML Engineer	\$23,000	\$40,000	\$69,000	\$56,000
Integration Engineer (EDI/FHIR)	\$14,000	\$25,000	\$44,000	\$35,000
DevOps / SRE Engineer	\$21,000	\$32,000	\$57,000	\$44,800
Security Engineer	\$17,000	\$32,000	\$52,000	\$44,800
QA Engineer	\$9,000	\$17,000	\$29,000	\$23,800
Data Engineer	\$14,000	\$28,000	\$48,000	\$39,200
Compliance Consultant (fractional)	\$18,000	\$36,000	\$60,000	\$50,400

18. Milestones — Full Detail

#	Milestone	Dur.	Key deliverables
M0	Discovery	2 wk	Requirements, payer/API/security research
M1	Architecture + Security	3 wk	Azure, Entra, networking, tenant model, CI/CD
M2	Multi-tenant core	4 wk	Tenants, RBAC, orgs, users, audit, admin
M3	Eligibility + Benefits	4 wk	270/271, integration framework, payer config
M4	Claim Status	3 wk	276/277, timeline, status agent
M5	Prior Authorization	4 wk	278/FHIR, PA workflow, documents
M6	Denial Intelligence	4 wk	835/CARC/RARC, analysis, recovery
M7	Appeals	3 wk	Draft, human approval, tracking
M8	Payer Intelligence	4 wk	Payer DB, RAG, recommendation engine
M9	AI Workforce	3 wk	Agent orchestration, registry, audit
M10	Security / HIPAA hardening	4 wk	Risk assessment, pen test, PHI controls
M11	Pilot	4 wk	Design partner, approved data, metrics
M12	Production readiness	4 wk	HA, DR, monitoring, runbooks, release

Total realistic ≈ 46 weeks. MVP-capable by ~M8–M9; pilot ~M11; production ~M12.

19. 12-Month & 24-Month Budgets

19.1 12-month budget

Category	USD
Team (Balanced 8, loaded)	\$414,400
Azure (MVP→Medium)	\$51,318
AI runtime	\$636,210
AI dev subscriptions	\$13,200
Security + SOC2 Type I	\$138,000
Legal + testing + clearinghouse	\$45,000
TOTAL (incl. 15% contingency)	\$1,492,847

19.2 24-month budget

Category	USD
Team (scale to production)	\$1,038,800
Azure (Medium × 24 mo)	\$172,224
AI runtime (Medium × 24 mo)	\$2,312,400
AI dev subscriptions × 24 mo	\$26,400
Security + SOC2 Type II	\$162,000
Legal + testing × 2	\$90,000
TOTAL (incl. 15% contingency)	\$4,450,298

20. Commercial Cost Models

Model	Cost structure	Gross-margin target
Provider-only	Azure + AI (variable/claim) + security (fixed) + clearinghouse (per txn)	70–80% at scale
Payer-only	Higher security/compliance + enterprise integration + human oversight	60–75%
Combined	Shared platform (amortized build) + tenant isolation + both-side integration	65–78% blended

Illustrative gross-margin structures; validate through pricing pilots. The combined model has the strongest LTV via governed cross-side network effects.

20.1 Build / Buy / Partner

Capability	Decision	Rationale
Payer connectivity	BUY/PARTNER	Clearinghouse (Stedi/Optum/Availity)
Denial/appeal/estimation AI	BUILD	Proprietary value
Payer Intelligence + Orchestration	BUILD	The moat
AI voice	BUY/INTEGRATE	Voice provider under BAA
Coverage criteria (InterQual/MCG)	PARTNER/LICENSE	License-gated
Payer-side adjudication	PARTNER/REGULATED	Decision-support only; legal review

21. Disaster Recovery, Continuity & Governance

21.1 RPO / RTO

Tier	RPO	RTO	Measures
MVP	≤ 24h	≤ 8h	Backups, PITR
Production	≤ 1h	≤ 4h	Geo-redundancy, failover, DR testing
Enterprise	≤ 15m	≤ 1h	Active-passive multi-region, ransomware recovery

21.2 Continuity responses

Scenario	Response
AI provider outage	Fallback: primary → secondary model → rules → human
Clearinghouse/payer API outage	Alternate channel; queue & retry; portal/voice
Azure / DB outage	Failover; restore from PITR/geo-backup
Cyberattack / ransomware	IR plan; isolation; immutable backups; restore

21.3 AI agent governance

AI agents are privileged software identities: agent identity → scoped tool/data/tenant/action permissions → audit. High-risk actions (production deploy, DB/IAM/security/PHI changes, deletions, and any adjudication/medical-necessity/coverage/financial decision) require human approval.

22. Connectivity Strategy

The orchestration engine follows one principle — **don't call when you can connect digitally** — selecting the cheapest sufficient channel per payer and workflow. Voice is a last resort after EDI, API, FHIR and portal; humans handle exceptions.

Workflow	Digital-first path	Fallback
Eligibility	270/271 → API → FHIR (CoverageEligibility*)	Portal → Voice → Human
Claim Status	276/277 → API → FHIR (Claim/ClaimResponse)	Portal → Voice → Human
Prior Auth	FHIR PAS / X12 278 / payer API	Portal → Voice → Human
Denials	835 / 277CA (electronic — no call)	Human review of exceptions
Appeals	Payer-specific: API / portal / upload	Fax → Mail → Human

Each execution captures outcome + evidence + confidence, schedules the next action, and escalates to a human below a confidence threshold. This right-channel routing is the core cost lever: it keeps the expensive channels (voice, human) rare.

23. Security Controls & PHI Data Flow

Domain	Controls
Identity	Entra ID · MFA · Conditional Access · PIM/JIT · Managed Identity
Network	Hub-spoke · Private Link/Endpoints · Firewall · WAF · DDoS · no public DB
Data	TLS 1.2+ · AES-256 at rest · Key Vault/CMK · HSM (where justified) · RLS
AI	AI Gateway: PHI detection → policy → redaction → approved model → validation → audit
Monitoring	Defender for Cloud · Sentinel SIEM · Monitor/App Insights · immutable audit
SDLC	SAST/DAST/SCA/secret/container/IaC scanning · protected main · PR review
Governance	Least privilege (humans + AI agents); high-risk actions human-approved

23.1 PHI data flow (AI Gateway)

Step	Action
1 Ingress	Authenticated request (Entra) over TLS
2 PHI detection	Classify & detect PHI in payload
3 Policy engine	Apply tenant + PHI policy; choose model + redaction
4 Redaction	Tokenize/redact where appropriate before the model call
5 Approved model	Azure OpenAI (BAA) for the PHI path; provider-agnostic gateway
6 Validation	Validate output; confidence scoring
7 Audit	Immutable log: input refs, decision, model, confidence, approver

Development data rule

Development uses synthetic / de-identified data only. Developer AI tools (Claude Max, ChatGPT subscriptions) are never production PHI infrastructure — production PHI uses only approved services under appropriate contractual/BAA arrangements.

24. Risks

- **Clearinghouse pricing opacity** — capture Stedi cents; start on the free tier; keep multi-vendor.
- **AI-decision regulation** (CMS-4201-F, CA SB 1120) — human-in-the-loop; decision-support only.
- **PHI/security** — Zero Trust, encryption, private networking, SIEM, incident-response readiness.
- **Tenant leakage** — RLS + app authorization + isolation testing.
- **AI cost overrun** — model routing, caching, batch, per-claim budget guardrails.
- **Timeline/talent** — AI productivity uplift, realistic plan, 15% contingency.

25. Final Recommendation

Build the common platform from day one; launch Provider/RCM first; add Payer capabilities after validation. The research supports this: provider-side connectivity is solved by clearinghouses (a near-\$0 MVP is feasible on free tiers), the durable moat is the Payer-Intelligence + Orchestration layer, and payer-side is a large adjacent opportunity best entered as decision-support. Provider-first minimizes regulatory complexity, capital and sales-cycle risk while proving the shared foundation.

Investor summary

MVP ≈ \$383,068; 12-month ≈ \$1,492,847; 24-month ≈ \$4,450,298; avg monthly burn (yr 1) ≈ \$124,404. Team: 5 FTE (lean) → 8 (balanced) → 12–15 (production), India-based, AI-augmented, 5 × Claude Max 20x. Azure \$1,377→\$38,694/mo across scale. HIPAA-aligned, SOC 2 on the roadmap. All figures are MODEL ASSUMPTIONS — verify live pricing before board commitment.

Appendix. Pricing Sources

Item	Confidence
Anthropic Claude Max plan	High
Anthropic Claude Team plan	High
Anthropic Claude Enterprise billing	Med-High
Claude Code with Pro/Max	High
OpenAI ChatGPT Business pricing	High
OpenAI ChatGPT Enterprise (quote-only)	High
Anthropic API pricing (secondary-corroborated)	Med
OpenAI API pricing	Med
Azure OpenAI Service pricing + HIPAA BAA	Med
Azure OpenAI HIPAA eligibility (Microsoft Learn)	High
Azure AI Search pricing	Med
Azure Container Apps pricing	High
Azure Database for PostgreSQL Flexible Server pricing	Med
Azure Front Door pricing	High
Azure API Management pricing	Med
Azure Service Bus pricing	Med
Azure Cache for Redis pricing	Med
Azure Blob Storage pricing	High
Azure Key Vault / Managed HSM pricing	Med
Microsoft Defender for Cloud pricing	High
Microsoft Sentinel pricing	High
Azure Monitor / Log Analytics pricing	High
Azure Firewall pricing	High
Azure Bastion pricing	Med
Azure Container Registry pricing	High
Azure Bandwidth pricing	High
Microsoft HIPAA BAA (Product Terms/DPA)	High
Azure AI Document Intelligence pricing	Med
Voice-AI stack per-minute (blended estimate)	Est
Stedi clearinghouse pricing (free tier + PAYG)	High
SOC 2 cost benchmarks (Secureframe/Drata)	Med-High
HITRUST CSF cost benchmarks (Thoropass/Sprinto)	Med
HIPAA risk assessment cost (Secureframe)	Med
Penetration test cost (Bright Defense)	Med-High
Compliance automation tooling cost	Med-High
vCISO / fractional CISO cost	Med-High
India developer salary benchmarks (Glassdoor/6figr/Levels)	Med

Full URLs are hyperlinked in the workbook SOURCES sheet. Primary hubs: Azure pricing pages / Microsoft Learn, Anthropic & OpenAI official pricing, CMS/HHS/HL7/X12/CAQH, and named benchmark sources for security/salary ranges.